

Models for each category

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON



Maarten Van den Broeck

Content Developer at DataCamp

Four categories

```
print(fish["species"].unique())
```

```
array(['Bream', 'Roach', 'Perch', 'Pike'], dtype=object)
```

Splitting the dataset

```
bream = fish[fish["species"] == "Bream"]  
perch = fish[fish["species"] == "Perch"]  
pike = fish[fish["species"] == "Pike"]  
roach = fish[fish["species"] == "Roach"]
```

Four models

```
mdl_bream = ols("mass_g ~ length_cm", data=bream).fit()  
print(mdl_bream.params)
```

```
Intercept    -1035.3476  
length_cm      54.5500
```

```
perch = fish[fish["species"] == "Perch"]  
print(mdl_perch.params)
```

```
Intercept    -619.1751  
length_cm     38.9115
```

```
mdl_pike = ols("mass_g ~ length_cm", data=pike).fit()  
print(mdl_pike.params)
```

```
Intercept    -1540.8243  
length_cm     53.1949
```

```
mdl_roach = ols("mass_g ~ length_cm", data=roach).fit()  
print(mdl_roach.params)
```

```
Intercept    -329.3762  
length_cm     23.3193
```

Explanatory data

```
explanatory_data = pd.DataFrame(  
    {"length_cm": np.arange(5, 61, 5)})  
  
print(explanatory_data)
```

	length_cm
0	5
1	10
2	15
3	20
4	25
5	30
6	35
7	40
8	45
9	50
10	55
11	60

Making predictions

```
prediction_data_bream = explanatory_data.assign(  
    mass_g = mdl_bream.predict(explanatory_data),  
    species = "Bream")
```

```
prediction_data_perch = explanatory_data.assign(  
    mass_g = mdl_perch.predict(explanatory_data),  
    species = "Perch")
```

```
prediction_data_pike = explanatory_data.assign(  
    mass_g = mdl_pike.predict(explanatory_data),  
    species = "Pike")
```

```
prediction_data_roach = explanatory_data.assign(  
    mass_g = mdl_roach.predict(explanatory_data),  
    species = "Roach")
```


Concatenating predictions

```
prediction_data = pd.concat([prediction_data_bream,  
                             prediction_data_roach,  
                             prediction_data_perch,  
                             prediction_data_pike])
```

	length_cm	mass_g	species
0	5	-762.597660	Bream
1	10	-489.847756	Bream
2	15	-217.097851	Bream
3	20	55.652054	Bream
4	25	328.401958	Bream
5	30	601.151863	Bream
...			
3	20	-476.926955	Pike
4	25	-210.952626	Pike
5	30	55.021703	Pike
6	35	320.996032	Pike
7	40	586.970362	Pike
8	45	852.944691	Pike
9	50	1118.919020	Pike
10	55	1384.893349	Pike
11	60	1650.867679	Pike


Visualizing predictions

```
sns.lmplot(x="length_cm",  
           y="mass_g",  
           data=fish,  
           hue="species",  
           ci=None)  
  
plt.show()
```



Adding in your predictions

```
sns.lmplot(x="length_cm",  
          y="mass_g",  
          data=fish,  
          hue="species",  
          ci=None)  
  
sns.scatterplot(x="length_cm",  
               y="mass_g",  
               data=prediction_data,  
               hue="species",  
               ci=None,  
               legend=False)  
  
plt.show()
```



Coefficient of determination

```
mdl_fish = ols("mass_g ~ length_cm + species",  
               data=fish).fit()
```

```
print(mdl_fish.rsquared_adj)
```

0.917

```
print(mdl_bream.rsquared_adj)
```

0.874

```
print(mdl_perch.rsquared_adj)
```

0.917

```
print(mdl_pike.rsquared_adj)
```

0.941

```
print(mdl_roach.rsquared_adj)
```

0.815

Residual standard error

```
print(np.sqrt mdl_fish.mse_resid))
```

103

```
print(np.sqrt mdl_bream.mse_resid))
```

74.2

```
print(np.sqrt mdl_perch.mse_resid))
```

100

```
print(np.sqrt mdl_pike.mse_resid))
```

120

```
print(np.sqrt mdl_roach.mse_resid))
```

38.2

Let's practice!

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON

One model with an interaction

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON



Maarten Van den Broeck

Content Developer at DataCamp

What is an interaction?

In the fish dataset

- Different fish species have different mass to length ratios.
- The effect of length on the expected mass is different for different species.

More generally

The effect of one explanatory variable on the expected response changes depending on the value of another explanatory variable.

Specifying interactions

No interactions

```
response ~ explntry1 + explntry2
```

With interactions (implicit)

```
response_var ~ explntry1 * explntry2
```

With interactions (explicit)

```
response ~ explntry1 + explntry2 + explntry1:explntry2
```

No interactions

```
mass_g ~ length_cm + species
```

With interactions (implicit)

```
mass_g ~ length_cm * species
```

With interactions (explicit)

```
mass_g ~ length_cm + species + length_cm:species
```

Running the model

```
mdl_mass_vs_both = ols("mass_g ~ length_cm * species", data=fish).fit()

print(mdl_mass_vs_both.params)
```

```
Intercept                -1035.3476
species[T.Perch]           416.1725
species[T.Pike]           -505.4767
species[T.Roach]           705.9714
length_cm                 54.5500
length_cm:species[T.Perch] -15.6385
length_cm:species[T.Pike]  -1.3551
length_cm:species[T.Roach] -31.2307
```


Easier to understand coefficients

```
mdl_mass_vs_both_inter = ols("mass_g ~ species + species:length_cm + 0", data=fish).fit()

print(mdl_mass_vs_both_inter.params)
```

```
species[Bream]          -1035.3476
species[Perch]          -619.1751
species[Pike]           -1540.8243
species[Roach]          -329.3762
species[Bream]:length_cm  54.5500
species[Perch]:length_cm  38.9115
species[Pike]:length_cm  53.1949
species[Roach]:length_cm  23.3193
```

Familiar numbers

```
print mdl_mass_vs_both_inter.params
```

```
species[Bream]          -1035.3476
species[Perch]          -619.1751
species[Pike]           -1540.8243
species[Roach]          -329.3762
species[Bream]:length_cm 54.5500
species[Perch]:length_cm 38.9115
species[Pike]:length_cm  53.1949
species[Roach]:length_cm 23.3193
```

```
print mdl_bream.params
```

```
Intercept    -1035.3476
length_cm     54.5500
```

Let's practice!

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON

Making predictions with interactions

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON



Maarten Van den Broeck

Content Developer at DataCamp

The model with the interaction

```
mdl_mass_vs_both_inter = ols("mass_g ~ species + species:length_cm + 0",  
                              data=fish).fit()  
  
print(mdl_mass_vs_both_inter.params)
```

```
species[Bream]          -1035.3476  
species[Perch]          -619.1751  
species[Pike]           -1540.8243  
species[Roach]          -329.3762  
species[Bream]:length_cm  54.5500  
species[Perch]:length_cm  38.9115  
species[Pike]:length_cm  53.1949  
species[Roach]:length_cm  23.3193
```

The prediction flow

```
from itertools import product
```

```
length_cm = np.arange(5, 61, 5)
```

```
species = fish["species"].unique()
```

```
p = product(length_cm, species)
```

```
explanatory_data = pd.DataFrame(p,  
                                columns=["length_cm",  
                                         "species"])
```


```
prediction_data = explanatory_data.assign(  
    mass_g = mdl_mass_vs_both_inter.predict(explanatory_data))
```

```
print(prediction_data)
```

	length_cm	species	mass_g
0	5	Bream	-762.5977
1	5	Roach	-212.7799
2	5	Perch	-424.6178
3	5	Pike	-1274.8499
4	10	Bream	-489.8478
5	10	Roach	-96.1836
6	10	Perch	-230.0604
7	10	Pike	-1008.8756
8	15	Bream	-217.0979
...			
40	55	Bream	1964.9014
41	55	Roach	953.1833
42	55	Perch	1520.9556
43	55	Pike	1384.8933
44	60	Bream	2237.6513
45	60	Roach	1069.7796
46	60	Perch	1715.5129
47	60	Pike	1650.8677

Visualizing the predictions

```
sns.lmplot(x="length_cm",  
           y="mass_g",  
           data=fish,  
           hue="species",  
           ci=None)  
  
sns.scatterplot(x="length_cm",  
                y="mass_g",  
                data=prediction_data,  
                hue="species")  
  
plt.show()
```



Manually calculating the predictions

```
coeffs = mdl_mass_vs_both_inter.params
```

```
species[Bream]          -1035.3476
species[Perch]           -619.1751
species[Pike]            -1540.8243
species[Roach]           -329.3762
species[Bream]:length_cm  54.5500
species[Perch]:length_cm  38.9115
species[Pike]:length_cm  53.1949
species[Roach]:length_cm  23.3193
```

```
ic_bream, ic_perch, ic_pike, ic_roach,  
slope_bream, slope_perch, slope_pike, slope_roach = coeffs
```


Manually calculating the predictions

```
conditions = [  
    explanatory_data["species"] == "Bream",  
    explanatory_data["species"] == "Perch",  
    explanatory_data["species"] == "Pike",  
    explanatory_data["species"] == "Roach"  
]  
  
ic_choices = [ic_bream, ic_perch, ic_pike, ic_roach]  
intercept = np.select(conditions, ic_choices)  
  
slope_choices = [slope_bream, slope_perch, slope_pike, slope_roach]  
slope = np.select(conditions, slope_choices)
```

Manually calculating the predictions

```
prediction_data = explanatory_data.assign(  
    mass_g = intercept + slope * explanatory_data["length_cm"])  
  
print(prediction_data)
```

```
prediction_data = explanatory_data.assign(  
    mass_g = mdl_mass_vs_both_inter.predict(explanatory_data))  
  
print(prediction_data)
```

	length_cm	species	mass_g
0	5	Bream	-762.5977
1	5	Roach	-212.7799
2	5	Perch	-424.6178
3	5	Pike	-1274.8499
4	10	Bream	-489.8478
5	10	Roach	-96.1836
...			
43	55	Pike	1384.8933
44	60	Bream	2237.6513
45	60	Roach	1069.7796
46	60	Perch	1715.5129
47	60	Pike	1650.8677

	length_cm	species	mass_g
0	5	Bream	-762.5977
1	5	Roach	-212.7799
2	5	Perch	-424.6178
3	5	Pike	-1274.8499
4	10	Bream	-489.8478
5	10	Roach	-96.1836
...			
43	55	Pike	1384.8933
44	60	Bream	2237.6513
45	60	Roach	1069.7796
46	60	Perch	1715.5129
47	60	Pike	1650.8677

Let's practice!

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON

Simpson's Paradox

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON



Maarten Van den Broeck

Content Developer at DataCamp

A most ingenious paradox!

Simpson's Paradox occurs when the trend of a model on the whole dataset is very different from the trends shown by models on subsets of the dataset.

trend = slope coefficient

Synthetic Simpson data

x	y	group
62.24344	70.60840	D
52.33499	14.70577	B
56.36795	46.39554	C
66.80395	66.17487	D
66.53605	89.24658	E
62.38129	91.45260	E

- 5 groups of data, labeled "A" to "E"

¹ https://www.rdocumentation.org/packages/datasauRus/topics/simpsons_paradox

Linear regressions

Whole dataset

```
mdl_whole = ols("y ~ x",  
                data=simpsons_paradox).fit()  
  
print(mdl_whole.params)
```

Intercept	-38.554
x	1.751


By group

```
mdl_by_group = ols("y ~ group + group:x + 0",  
                   data = simpsons_paradox).fit()  
  
print(mdl_by_group.params)
```

groupA	groupB	groupC	groupD	groupE
32.5051	67.3886	99.6333	132.3932	123.8242
groupA:x	groupB:x	groupC:x	groupD:x	groupE:x
-0.6266	-1.0105	-0.9940	-0.9908	-0.5364


Plotting the whole dataset

```
sns.regplot(x="x",  
            y="y",  
            data=simpsons_paradox,  
            ci=None)
```



Plotting by group

```
sns.lmplot(x="x",  
          y="y",  
          data=simpsons_paradox,  
          hue="group",  
          ci=None)
```



Reconciling the difference

Good advice

If possible, try to plot the dataset.



Common advice

You can't choose the best model in general – it depends on the dataset and the question you are trying to answer.



More good advice

Articulate a question before you start modeling.

Test score example



Infectious disease example



Reconciling the difference

- Usually (but not always) the grouped model contains more insight.
- Are you missing explanatory variables?
- Context is important.

Simpson's paradox in real datasets

- The paradox is usually less obvious.
- You may see a zero slope rather than a complete change in direction.
- It may not appear in every group.

Let's practice!

INTERMEDIATE REGRESSION WITH STATSMODELS IN PYTHON